


# Recap

- **Linear Regression (Simple + Multiple)**
- **Objective of linear regression**
- **Assumptions of linear regression**
- **Building a linear regression model in python: sklearn and statsmodels**
- **Interpreting simple + multiple linear regression**
- **Evaluation metrics for regression: MAE, MSE, RMSE**
- **Train/test split in python**

**Week 6: Data Science Part-Time Course**

# **Polynomial Regression, Regularization and Resampling**

Dami Lasisi



**Linear Regression:**

$$\text{Tip} = b_0 + b_1 \text{Total restaurant bill} + \varepsilon$$

**Polynomial Regression:**

$$\text{Tip} = b_0 + b_1 \text{Total restaurant bill} + b_2 \text{Total restaurant bill}^2 + \varepsilon$$

Higher order polynomials allow us to produce extremely non-linear curves, but don't get too crazy because higher orders can create very weird shapes

# Bias and Variance






Fig. 1 Graphical illustration of bias and variance.

SOURCE: Understanding the Bias-Variance Tradeoff, by Scott Fortmann-Roe.

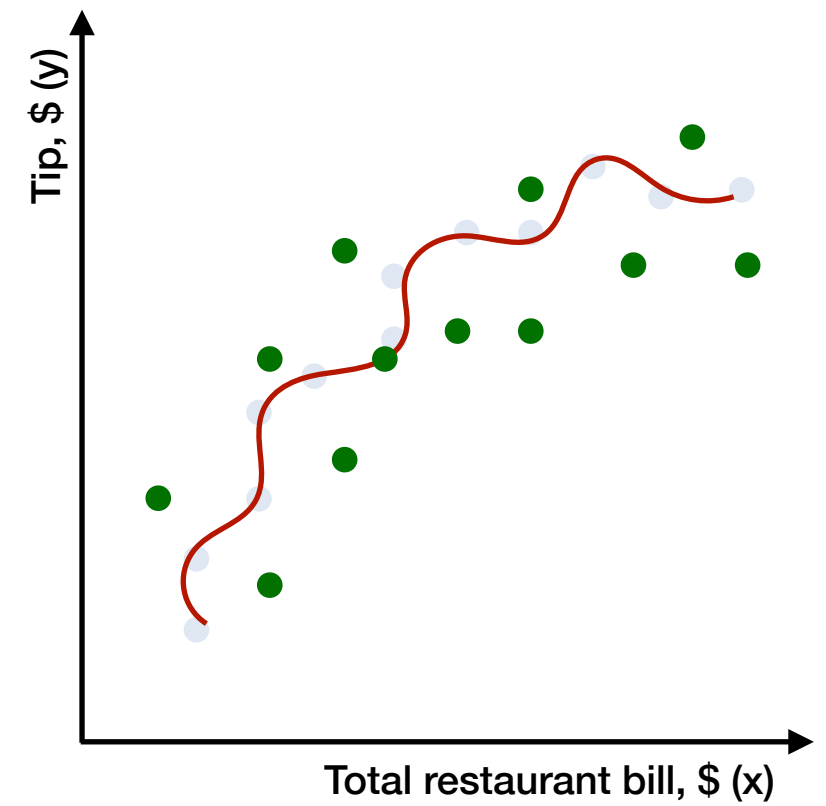
# Bias and Variance



High bias - underfitting




Sweet spot




High variance - overfitting

# Bias and Variance



# Bias Variance Tradeoff



# Dealing with Overfitting: Regularization

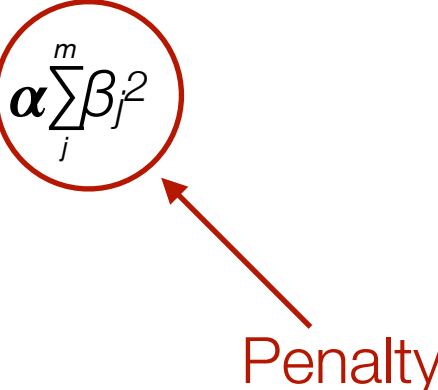
- ▶ Imposing penalties on complex models
- ▶ Regularization is a method for "constraining" or "regularizing" the size of the coefficients, thus "shrinking" them toward zero
- ▶ It reduces model variance and thus minimizes overfitting
- ▶ If the model is too complex, it tends to reduce variance more than it increases bias, resulting in a model that is more likely to generalize
- ▶ **Regularization techniques:** Ridge regression, Lasso regression, Elastic net regression



# Ridge Regression

- ▶ Reduces model complexity
- ▶ Quantifies overfitting through measure of magnitude of coefficients
- ▶ Includes a tuning parameter ( $\alpha$ ) that will decrease the magnitude of the coefficients to approach zero
- ▶ Shrinks parameters so it is mostly used to prevent multicollinearity
- ▶ Uses L2 regularization technique

▶ **Minimizes:**  $\sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha \sum_j^m \beta_j^2$

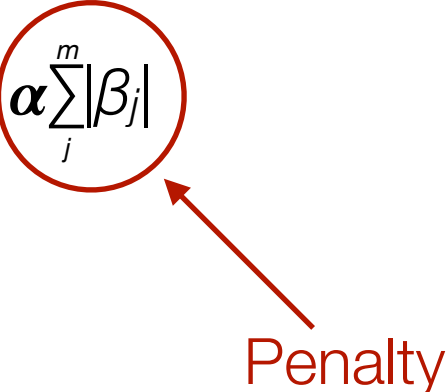


Penalty

# Lasso Regression

- ▶ Least **A**bsolute **S**hrinkage **S**elector **O**perator
- ▶ Built-in feature selection: uses tuning parameter ( $\alpha$ ) to decrease the magnitude of the coefficients (sometimes to absolute zeros)
- ▶ Generally used when we have more numbers of features
- ▶ Uses L1 regularization technique

▶ **Minimizes:**  $\sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha \sum_j^m |\beta_j|$

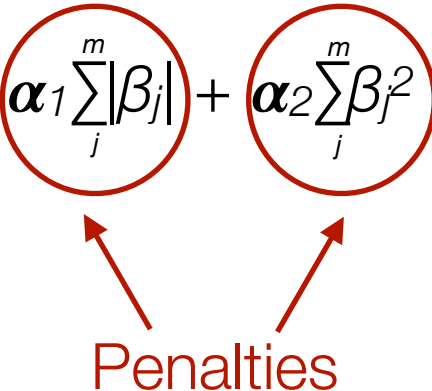


Penalty

# Elastic Net Regression

- ▶ A hybrid of lasso and ridge regression
- ▶ Works well when we have a large set of features
- ▶ Uses both L1 and L2 regularization techniques

▶ **Minimizes:**  $\sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha_1 \sum_j^m |\beta_j| + \alpha_2 \sum_j^m \beta_j^2$



Penalties

# Implementing regularization in python

```
In [ ]: from sklearn.linear_model import Ridge, Lasso, ElasticNet
```

```
In [ ]: X = data[features]
        y = data[target]

        #Ridge
        rr = Ridge(alpha=2.0)
        rr.fit(X,y)
        rr.predict(X)

        #Lasso
        lr = Lasso(alpha=2.0)
        lr.fit(X,y)
        lr.predict(X)

        #Elastic Net
        er = ElasticNet(alpha=2.0, l1_ratio=0.5)
        er.fit(X,y)
        er.predict(X)
```

# Applying regularization

- ▶ Standardizing the features
  - To avoid penalizing the features simply because of their scale
  - To avoid penalizing the intercept, **which wouldn't make sense**

```
In [ ]: from sklearn.preprocessing import StandardScaler
```

```
In [ ]: X = data[features]
        y = data[target]

        ss = StandardScaler()
        X_stand = ss.fit_transform(X)
```

- ▶ When to choose lasso regression or ridge regression
  - Lasso regression is preferred if we believe many features are irrelevant or if we prefer a sparse model.
  - Ridge can work particularly well if there is a high degree of multicollinearity in your model.
  - If model performance is your primary concern, it is best to try both using Elastic net regression

# Cross Validation

- ▶ Used to estimate the performance and skill of a machine learning model
- ▶ Reserves a portion of the dataset that is not included in training the model
- ▶ Steps:
  - Reserve a portion of the dataset (test set)
  - Train the model with the unreserved portion of the dataset (train set)
  - Use the reserved set to validate the model by measuring its performance

```
In [ ]: from sklearn import metrics
        from sklearn.model_selection import train_test_split
```

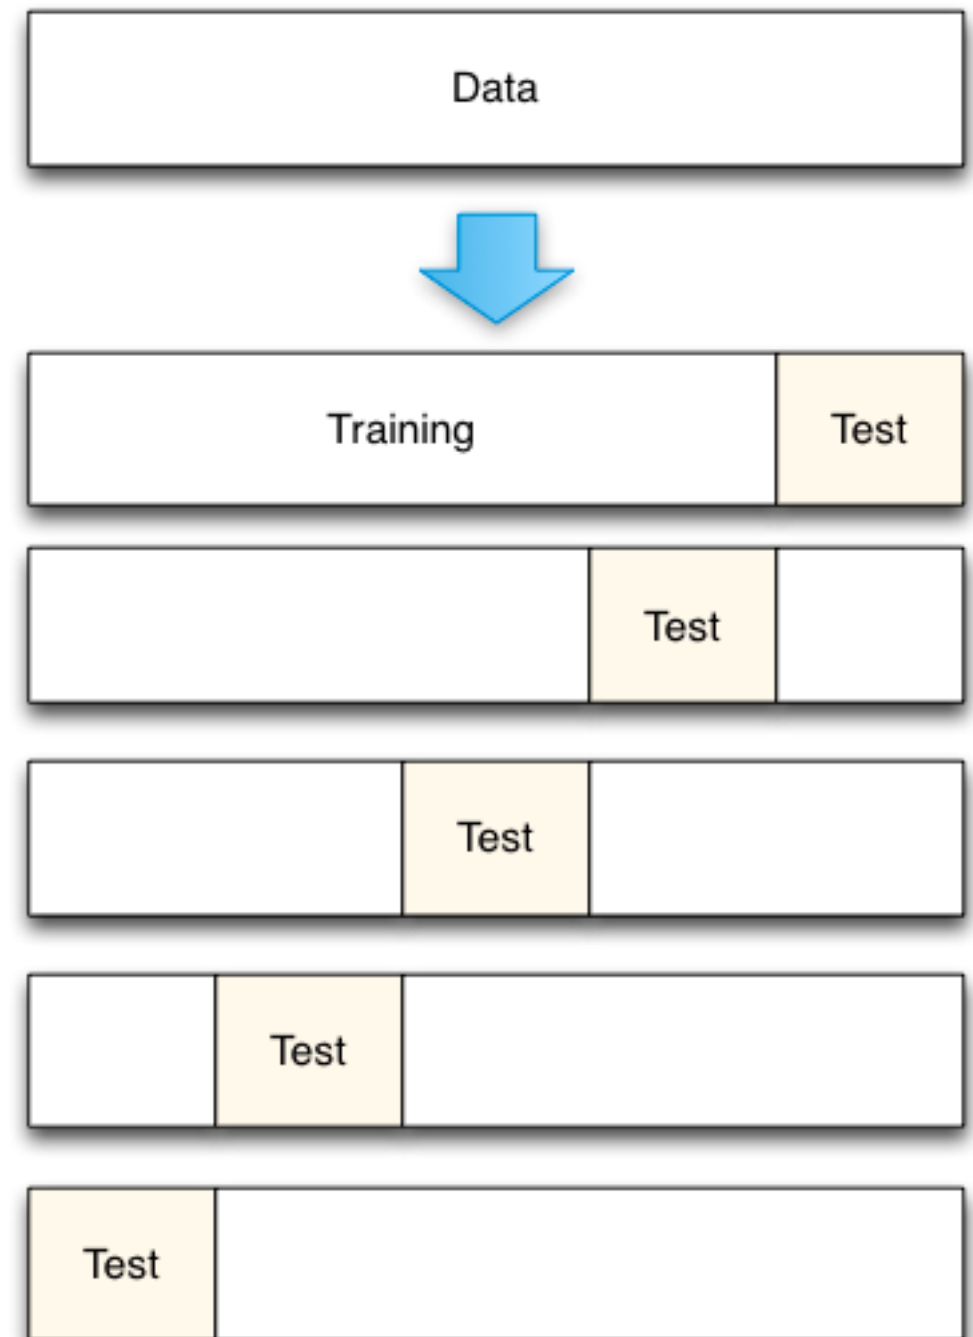
```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, random_state=123)

        linreg = LinearRegression()
        linreg.fit(X_train, y_train)

        y_pred = linreg.predict(X_test)
        print("MSE:", np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

# K-Fold Cross Validation

- ▶ Randomly shuffle the dataset
- ▶ Split the dataset into k groups
- ▶ For each group, train you model with k-1 groups in the dataset and then test with kth group
- ▶ Record the error score for each model built
- ▶ Take the average of your k recorded error scores. This is the cross-validation error that serves as your performance metric for the model.



# Leave-One-Out Cross Validation

- ▶ We train  $n-1$  observations in the dataset and then test on the remaining one
- ▶ This process iterates for each data point



# Three-Way Data Split

